

Changmin Choi (최창민)

cmchoi9901@gmail.com | 010-4403-4070

github.com/ryul99 | torch.vision

Education

- **Seoul National University, Computer Science and Engineering**, Mar 2017 – Feb 2024
 - Bacchus: System Administration Club (Dept. of CSE) – Server Administrator / Vice-chair

Careers

- **Bucketplace (Ohouse / 오늘의집)** (Jan 2025 – Present): Machine Learning Engineer
 - Developed a vector-search-driven semantic retrieval system (Tech Blog)
 - Fine-tuned BGE-M3 on proprietary data; introduced GOR loss and enabled MRL-style embeddings
 - Evaluated multiple serving stacks, recommended the best-fit framework
 - Built a continuous quality loop using LLM-based evaluation to drive iterative improvements
 - Delivered +2.45% user interactions and +2.82% CTR
 - Owned object detection for an image-based product search experience (Tech Blog)
 - Implemented pipelines with robust data validation; identified and fixed multiple data consistency bugs
 - Designed post-processing tailored to internal data characteristics to improve detection quality
 - Delivered +5.24% GMV and +11.78% commerce profit
- **FuriosaAI** (Feb 2024 – Dec 2024): Solution Architect
 - Improved end-to-end NPU performance via the Python SDK
 - Optimized the pipeline to overlap CPU preprocessing/postprocessing with NPU inference
 - Built and benchmarked TensorRT-LLM serving stack; measured latency, throughput, and power efficiency
- **Moreh** (Jul 2023 – Dec 2023): Part-time Intern (AI Compiler Team)
 - Measured and improved an auto-parallelized GPU training cluster
 - Instrumented training-time GPU throughput and surfaced issues in the throughput estimator
- **Hyperconnect** (Sep 2020 – Aug 2021): Machine Learning Engineer
 - Developed next-generation recommendation models
 - Identified limitations of a frequently retrained CatBoost-based recommender; implemented deep-learning alternatives
 - Built and operated recommendation dashboards
 - Built monitoring dashboards using Google BigQuery
 - Improved cross-section comparability between rule-based and ML models via calibrated scaling
- **MaumAI** (Aug 2019 – Sep 2020): AI Researcher
 - Video Super Resolution
 - GitHub repo: <https://github.com/ryul99/mmsr>
 - Extended an open-source baseline (e.g., perceptual loss) and fine-tuned it for internal datasets
 - Contributed upstream with bug fixes and memory-usage optimizations
- **Seoul National University Hospital (SNUH CMI)** (Apr 2019 – Jul 2019): Researcher & Server Administrator
 - DevEnv with Docker
 - GitHub repo: <https://github.com/ryul99/DevEnv-with-docker>
 - Delivered isolated dev environments via Docker while aligning UID inside/outside containers
 - Enforced per-user memory limits on shared servers using cgroups

Skills

- **Machine Learning:** vector search, LLM, PyTorch, paper review, ML/DL fundamentals
- **Languages:** Python (proficient), Shell Script, C++, SQL, C#, Java, etc.
- **Familiar with:** Linux, macOS, Windows, Docker, cloud services, etc.